

CS331 Tutorial 3


Java Thread Programming (Locking and CDS)

A. Sahu

Dept of Comp. Sc. & Engg.

Indian Institute of Technology Guwahati

Some slides are adopted from
“The Art of Multiprocessor
Programming”
by Maurice Herlihy & Nir Shavit



Outline

- Java Threading Examples
- Java Incrementing Counter: Multithreaded Version
- Locking in Java
- Concurrent Data Structure
 - List, Hash, Queue, Stack
- Simulation of PDES

Thread Programing In Java

Create Java Thread: Runnable Interface

```
public class MyThIR implements Runnable{  
    public void run() {  
        System.out.println("Child  
        thread is running...");  
    }  
}  
  
public static void main(String args[]) {  
    MyThIR m1=new MyThIR();  
    Thread t1 =new Thread(m1);  
    t1.start();  
}
```

Create Java Thread: Extending Java Thread

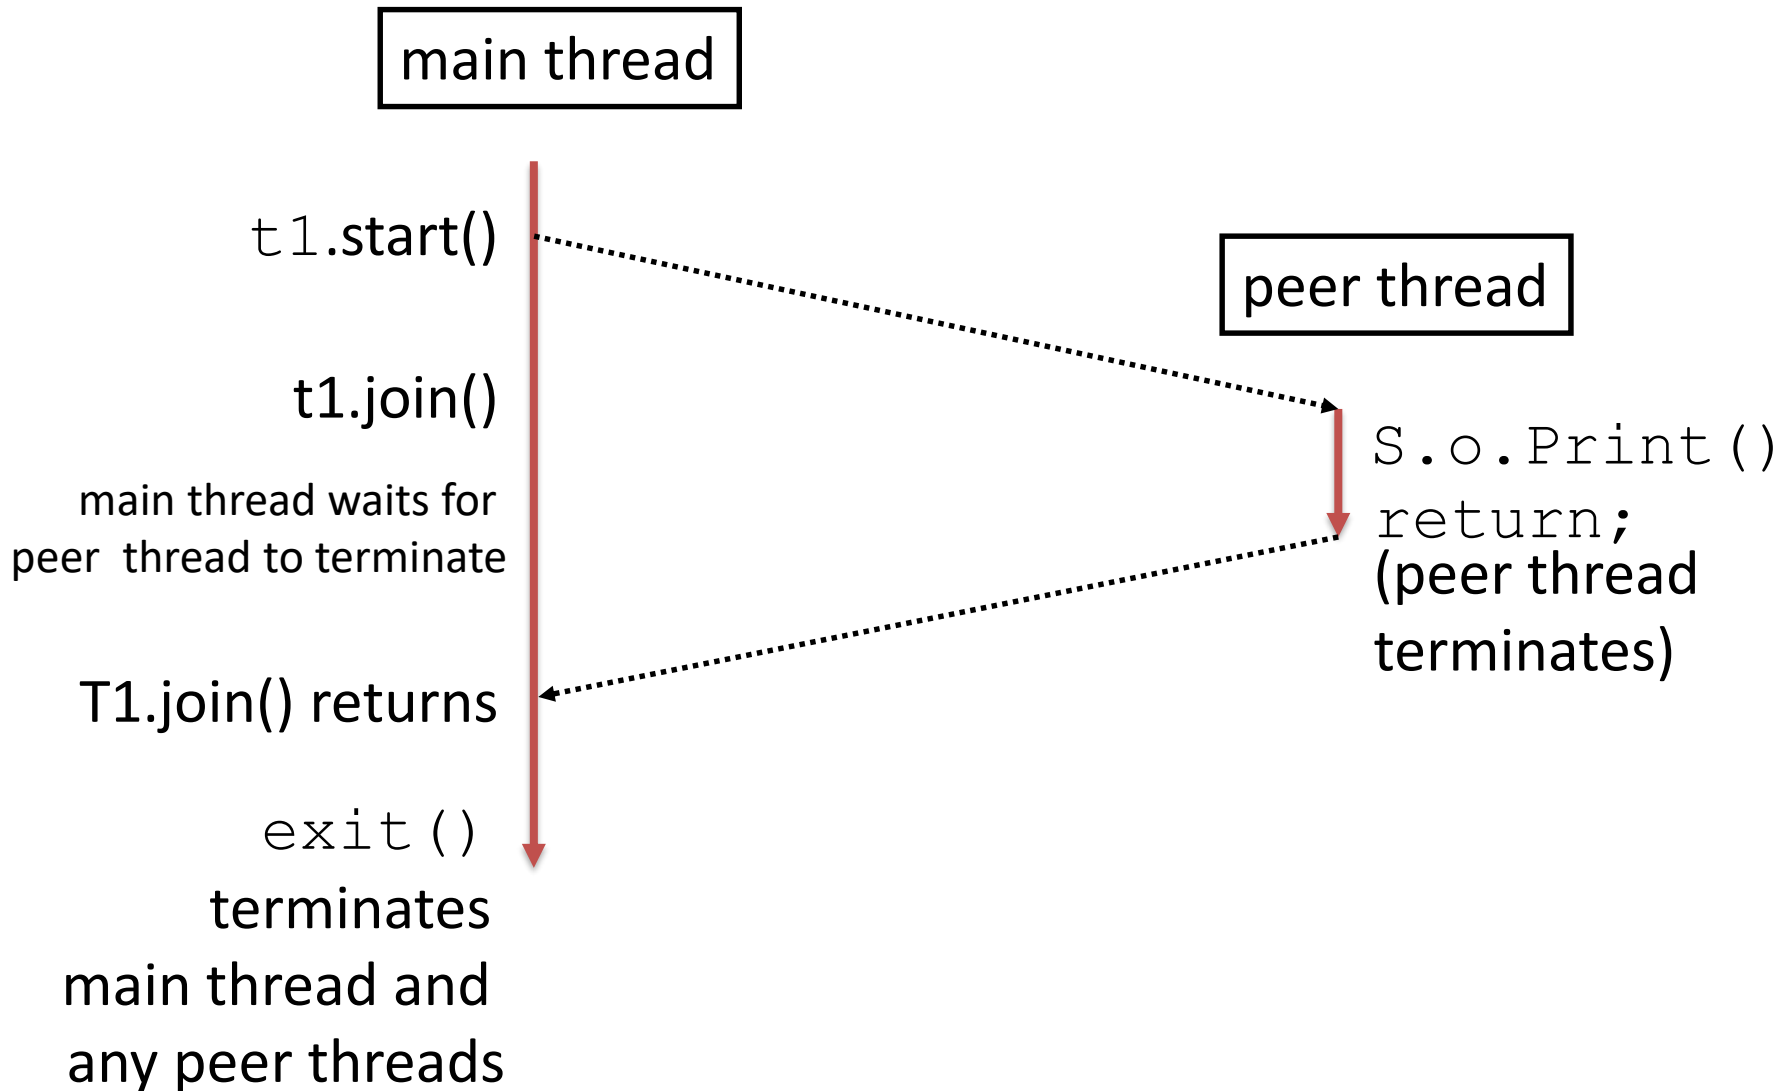
```
public class MyTh_EJT extends Thread{  
    public void run() {  
        System.out.println("Child  
        thread is running...");  
    }  
}  
  
public static void main(String args[]) {  
    MyTh_EJT t1 = new MyTh_EJT();  
    t1.start();  
}
```

Create Java Thread: Extending Java Thread

```
public class MyTh_EJT extends Thread{
    public void run() {
        System.out.println("Child
            thread is running...");
        return;
    }
}

public static void main(String args[]) {
    MyTh_EJT t1 = new MyTh_EJT();
    t1.start();
    t1.join();
}
```

Execution of Threaded “child thread”



VectorSum Serial

```
int A[VSize], B[VSize], C[VSize];  
void VectorSumSerial() {  
    for( int j=0; j<SIZE; j++)  
        A[j]=B[j]+C[j];  
}
```

Suppose Size=1000

0-249	250-499	500-749	750-999
-------	---------	---------	---------

T1

T2

T3

T4

VectorSum Serial

```
int A[VSize], B[VSize], C[VSize];  
void VectorSumSerial() {  
    for( int j=0; j<SIZE; j++)  
        A[j]=B[j]+C[j];  
}
```

- Independent
- Divide work into equal for each thread
- Work per thread: $\text{Size}/\text{numThread}$

VectorSum Parallel

```
void DoVectorSum(int TID) {  
    int j, SzPerthrd, LB, UB;  
    SzPerthrd=(VSize/NUM_THREADS);  
    LB= SzPerthrd*TID;UB=LB+SzPerthrd;  
  
    for (j=LB; j<UB; j++)  
        A[j]=B[j]+C[j];  
}
```

VectorSum Parallel : Java Threads

```
public class MyTh_EJT extends Thread{
    public void run(int TID) {
        DoVectorSum(TID)
        return;
    }
}

public static void main(String args[]) {
    int NumOfTasks = Integer.parseInt(args[0]);
    for ( int i=0; i < NumOfTasks; i++) {
        MyTh_EJT t1=new MyTh_EJT();
        t1.start(i);
    }
}
```

VectorSum Parallel : Java Threads

Complete java program will be
uploaded to MS Team

Four ways to implement a synchronized counter in Java

- **Suppose there is a Shared counter and every threads are attempting to manipulate**
 - 1. Synchronized Block**
 - 2. Atomic Variable**
 - 3. Concurrent Lock**
 - 4. Semaphore**

0: Simple Java Ctr: without Lock

```
import java.util.concurrent.ExecutorService;
```

```
import java.util.concurrent.Executors;
```

```
class Counter implements Runnable {
```

```
    private static int counter = 0;
```

```
    private static final int limit = 1000;
```

```
    private static final int threadPoolSize = 5;
```

```
    public static void main(String[] args) {
```

```
        ExecutorService ES = Executors.newFixedThreadPool(threadPoolSize);
```

```
        for (int i = 0; i < threadPoolSize; i++) { ES.submit(new Counter()); } 
```

```
        ES.shutdown();
```

```
    }
```

```
}
```

0: Simple Java Ctr: without Lock

```
@Override
public void run() {
    while (counter < limit) {
        increaseCounter();
    }
}

private void increaseCounter() {
    System.out.println(Thread.currentThread().getName() + " : " + counter);
    counter++;
}
}
```

1: Simple Java Ctr: with Sync Lock

```
import java.util.concurrent.ExecutorService;
import java.util.concurrent.Executors;

class Counter implements Runnable {
    private static int counter = 0;
    private static final int limit = 1000;
    private static final int threadPoolSize = 5;
    private static final Object lock = new Object();
    public static void main(String[] args) {
        ExecutorService ES = Executors.newFixedThreadPool(threadPoolSize);
        for (int i = 0; i < threadPoolSize; i++) { ES.submit(new Counter()); }
        ES.shutdown();
    }
}
```


1: Simple Java Ctr: with Sync Lock

```
@Override
public void run() {
    while (counter < limit) {
        increaseCounter();
    }
}
private void increaseCounter() {
    synchronized (lock) {
        System.out.println(Thread.currentThread().getName() + " : " + counter);
        counter++;
    }
}
}
```

2: Simple Java Ctr: with atomic Int

```
import java.util.concurrent.ExecutorService;
import java.util.concurrent.Executors;
import java.util.concurrent.atomic.AtomicInteger;
class Counter implements Runnable {
    private static AtomicInteger counter;
    private static final int limit = 1000;
    private static final int threadPoolSize = 5;
    public static void main(String[] args) {
        ExecutorService ES = Executors.newFixedThreadPool(threadPoolSize);
        for (int i = 0; i < threadPoolSize; i++) { ES.submit(new Counter()); }
        ES.shutdown();
    }
}
```

2: Simple Java Ctr: with atomic int

```
@Override
public void run() {
    while (counter < limit) {
        increaseCounter();
    }
}

private void increaseCounter() {
    System.out.println(Thread.currentThread().getName() + " : " + counter);
    counter.getAndIncrement();
}
}
```

3: Simple Java Ctr: Using concurrent lock

```
import java.util.concurrent.ExecutorService;
import java.util.concurrent.Executors;
import java.util.concurrent.locks.ReentrantLock;

class Counter implements Runnable {
    private ReentrantLock lock;
    private static final int limit = 1000;
    private static final int threadPoolSize = 5;

    public static void main(String[] args) {
        ReentrantLock sharedLock = new ReentrantLock();
        ExecutorService ES = Executors.newFixedThreadPool(threadPoolSize);
        for (int i = 0; i < threadPoolSize; i++) { ES.submit(new Counter(sharedLock)); }
        ES.shutdown();
    }
}
```

3: Simple Java Ctr: Using concurrent lock

```
public Counter(ReentrantLock lock) {  
    this.lock = lock;  
}  
@Override  
public void run() {  
    while (counter < limit) {  
        increaseCounter();  
    }  
}  
private void increaseCounter() {  
    lock.lock();  
    try { counter++; }  
    finally { lock.unlock(); }  
}
```

4: Simple Java Ctr: Using a Semaphore

```
import java.util.concurrent.ExecutorService;
import java.util.concurrent.Executors;
import java.util.concurrent.Semaphore;
class Counter implements Runnable {
    Semaphore semaphore;
    private static final int limit = 1000;
    private static final int threadPoolSize = 5;
    public static void main(String[] args) {
        Semaphore sharedSemaphore = new Semaphore(1);
        ExecutorService ES = Executors.newFixedThreadPool(threadPoolSize);
        for (int i = 0; i < threadPoolSize; i++) {
            ES.submit(new Counter(sharedSemaphore)); }
        ES.shutdown();
    }
}
```

4: Simple Java Ctr: Using a Semaphore

```
public Counter(Semaphore semaphore){
    this.semaphore = semaphore;
}
@Override
public void run() {
    while (counter < limit) {
        increaseCounter();
    }
}
private void increaseCounter() {
    try { semaphore.acquire(); counter++;
    } catch (InterruptedException e) { e.printStackTrace();
    } finally { semaphore.release();}
}
```

Basic Locking implementation in Java


Basic Locking implementation in Java

- Locking require BASIC, Atomic Instruction at OS level
- CAS : Compare and Set /Test and Increment
- TAS : Test and Set
- TTAS: Try, Test and Set
- Exponential Back up TAS

Synchronization Hierarchy ☺ ☺ ☺

- One == (used by) == > other
- **LL+SC ==> TAS/CAS/FAI/XCHG==>Lock/Unlock**
 - All **TAS/CAS/GAS/FAI/XCHG** do the same work
- Lock/Unlock == > Mutex //Mutex use L/UL
- Mutex == > Semaphore // Semaphore uses Mutex
 - Wait() and Signal()
- Semaphore == > Monitor //Monitor uses Semaphore
 - Many wait/Many Signal, Processes in Queue
 - Monitor : Another Abstract Type
 - *which use semaphore, mutex, conditions*

Many threads trying to acquire Mutex LOCK : Spin Lock



CAS: CompareAndSet

```
public abstract class Lock{
    private int value;
    public boolean synchronized
        CompareAndSet(int expected,
                      int update) {
        int prior = value;
        if (value==expected) {
            value = update; return true;
        }
        return false;
    }
}
```

Test-and-Set

- Boolean value
- Test-and-set (TAS)
 - Swap **true** with current value
 - Return value tells if prior value was **true** or **false**
- Can reset just by writing **false**
- TAS aka “getAndSet”

Test-and-Set

```
public class AtomicBoolean {  
    boolean value;  
  
    public synchronized boolean  
        getAndSet(boolean newValue) {  
        boolean prior = value;  
        value = newValue;  
        return prior;  
    }  
}
```

Test-and-set Lock : TASLock

```
class TASlock {  
    AtomicBoolean state =  
        new AtomicBoolean(false) ;  
  
    void lock() {  
        while (state.getAndSet(true)) {} //Spining  
    }  
  
    void unlock() {  
        state.set(false) ;  
    }  
}
```


Test-and-Set Locks

- Locking
 - Lock is free: value is false
 - Lock is taken: value is true
- Acquire lock by calling TAS
 - If result is false, you win
 - If result is true, you lose
- Release lock by writing false


Performance

- Experiment
 - n threads
 - Increment shared counter 1 million times
- How long should it take?
- How long does it take?

Graph



Mystery #1



Test-and-Test-and-Set Locks

- Lurking stage
 - Wait until lock “looks” free
 - Spin while read returns true (lock taken)
- Pouncing state
 - As soon as lock “looks” available
 - Read returns false (lock free)
 - Call TAS to acquire lock
 - If TAS loses, back to lurking

Test-and-test-and-set Lock

```
class TTASlock {  
    AtomicBoolean state =  
        new AtomicBoolean(false);  
  
    void lock() {  
        while (true) {  
            while (state.get()) {}  
            if (!state.getAndSet(true))  
                return;  
        }  
    }  
}
```

Test-and-test-and-set Lock

```
class TTASlock {  
    AtomicBoolean state =  
        new AtomicBoolean(false);  
  
    void lock() {  
        while (true) {  
            while (state.get()) {}  
            if (!state.getAndSet(true))  
                return;  
        }  
    }  
}
```

Wait until lock looks free

Test-and-test-and-set Lock

```
class TTASlock {  
    AtomicBoolean state =  
        new AtomicBoolean(false);
```

```
    void lock() {  
        while (true) {  
            while (state.get()) {}
```


```
            if (!state.getAndSet(true))  
                return;
```

```
        }  
    }  
}
```

Then try to
acquire it




Mystery #2




Solution: Introduce Delay

- If the lock looks free
 - But I fail to get it
- There must be contention
 - Better to back off than to collide again



Dynamic Example: Exponential Backoff



If I fail to get lock

- wait random duration before retry
- Each subsequent failure doubles expected wait

Exponential Backoff Lock

```
public class Backoff implements lock {  
    public void lock() {  
        int delay = MIN_DELAY;  
        while (true) {  
            while (state.get()) {}  
            if (!lock.getAndSet(true))  
                return;  
            sleep(random() % delay);  
            if (delay < MAX_DELAY)  
                delay = 2 * delay;  
        }  
    }  
}
```

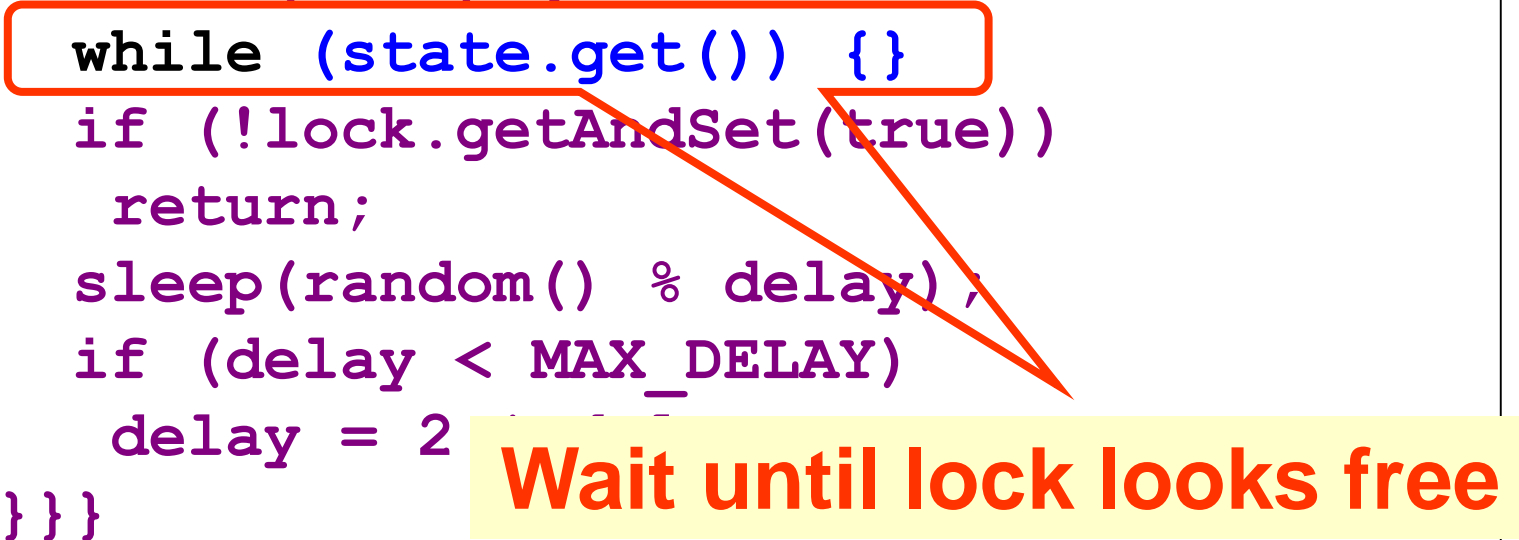
Exponential Backoff Lock

```
public class Backoff implements lock {  
    public void lock() {  
        int delay = MIN_DELAY;  
        while (true) {  
            while (state.get()) {}  
            if (!lock.getAndSet(true))  
                return;  
            sleep(random() % delay);  
            if (delay < MAX_DELAY)  
                delay = 2 * delay;  
        }  
    }  
}
```

Fix minimum delay

Exponential Backoff Lock

```
public class Backoff implements lock {  
    public void lock() {  
        int delay = MIN_DELAY;  
        while (true) {  
            while (state.get()) {}  
            if (!lock.getAndSet(true))  
                return;  
            sleep(random() % delay);  
            if (delay < MAX_DELAY)  
                delay = 2 * delay;  
        }  
    }  
}
```



Wait until lock looks free

Exponential Backoff Lock

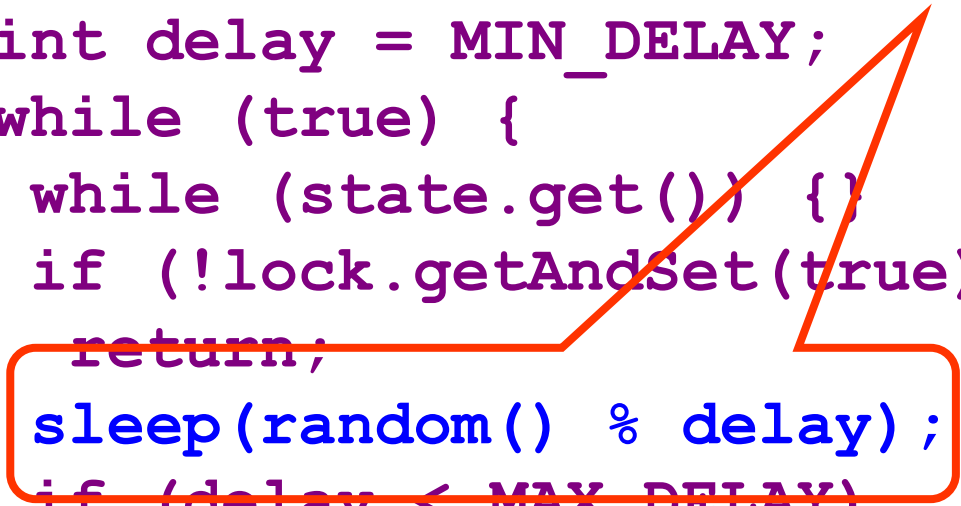
```
public class Backoff implements lock {  
    public void lock() {  
        int delay = MIN_DELAY;  
        while (true) {  
            while (state.get()) {}  
            if (!lock.getAndSet(true))  
                return;  
            sleep(random() % delay);  
            if (delay < MAX_DELAY)  
                delay = 2 * delay;  
        }  
    }  
}
```

If we win, return

Exponential Backoff Lock

```
public class Backoff implements Lock {  
    public  
        int delay = MIN_DELAY;  
        while (true) {  
            while (state.get()) {}  
            if (!lock.getAndSet(true))  
                return;  
            sleep(random() % delay);  
            if (delay < MAX_DELAY)  
                delay = 2 * delay;  
        }  
    }  
}
```

Back off for random duration




Exponential Backoff Lock

```
public class Backoff implements Lock {  
    public  
    int delay = MIN_DELAY;  
    while (true) {  
        while (state.get()) {}  
        if (!lock.getAndSet(true))  
            return;  
        sleep(random() % delay);  
        if (delay < MAX_DELAY)  
            delay = 2 * delay;  
    }  
}
```

Double max delay, within reason

Spin-Waiting Overhead



Backoff: Other Issues

- Good
 - Easy to implement
 - Beats TTAS lock
- Bad
 - Must choose parameters carefully
 - Not portable across platforms

Concurrent/Thread-Safe Data Structure

CDS or TS-DS

Data Structure as Object

- Data Stored on DS
 - List, Queue, Stack , Hash, Tree
- Concurrent access should be safe
 - All operation including read, write, modify, add, remove
- Safe : It should not alter the properties of the DS

Today: Concurrent Objects

- Adding threads should not lower throughput
 - Contention effects
 - Mostly fixed by Queue locks

Today: Concurrent Objects

- Adding threads should not lower throughput
 - Contention effects
 - Mostly fixed by Queue locks
- Should increase throughput
 - Not possible if inherently sequential
 - Surprising things are parallelizable

Coarse-Grained Synchronization

- Each method locks the object
 - Avoid contention using queue locks
 - Easy to reason about
 - In simple cases
- So, are we done?

Coarse-Grained Synchronization

- Sequential bottleneck
 - Threads “stand in line”
- Adding more threads
 - Does not improve throughput
 - Struggle to keep it from getting worse
- So why even use a multiprocessor?
 - Well, some apps inherently parallel ...

Example : Linked List of Student

- Linked List of Students, with Roll No, Marks
- Roll No is Key
- Mark is Data
- Many thread simultaneously working on the LinkedList Object
- Operation to be performed
 - Read mark of a student
 - **Update mark of a student**
 - **Add a student**
 - **Delete a student**
 - Calculate average

```
class node {  
    int Roll;  
    int Marks  
    node next;  
};
```

Design of Concurrent DS

- Introduce four “patterns”
 - Bag of tricks ...
 - Methods that work more than once ...
- For highly-concurrent objects
 - Concurrent access
 - More threads, more throughput

First: Fine-Grained Synchronization

- Instead of using a single lock ...
- Split object into
 - Independently-synchronized components
 - **Example: Hash : Modification**
- Methods conflict when they access
 - The same component ...
 - At the same time

Second: Optimistic Synchronization

- Search without locking ...
- If you find it, lock and check ...
 - OK: we are done
 - Oops: start over
- Evaluation
 - Usually cheaper than locking, but
 - Mistakes are expensive

Third: Lazy Synchronization

- Postpone hard work
- Removing components is tricky
 - Logical removal
 - Mark component to be deleted
 - Physical removal
 - Do what needs to be done

Fourth: Lock-Free Synchronization

- Don't use locks at all
 - Use `compareAndSet()` & relatives ...

Fourth: Lock-Free Synchronization

- Don't use locks at all
 - Use `compareAndSet()` & relatives ...
- Advantages
 - No Scheduler Assumptions/Support
- Disadvantages
 - Complex
 - Sometimes high overhead

Thanks