

CS331 Tutorial 4

Java : CDS

A. Sahu

Dept of Comp. Sc. & Engg.

Indian Institute of Technology Guwahati

Outline

- Four Way to implement Lock in Java
- Concurrent List
- Concurrent Hash
- Assignment-II
 - Use may use C-List and C-Hash

Four ways to implement a synchronized counter in Java

- **Suppose there is a Shared counter and every threads are attempting to manipulate**
 - 1. Synchronized Block**
 - 2. Atomic Variable**
 - 3. Concurrent Lock**
 - 4. Semaphore**

First: Fine-Grained Synchronization

- Instead of using a single lock ...
- Split object into
 - Independently-synchronized components
 - **Example: Hash : Modification**
- Methods conflict when they access
 - The same component ...
 - At the same time

Second: Optimistic Synchronization

- Search without locking ...
- If you find it, lock and check ...
 - OK: we are done
 - Oops: start over
- Evaluation
 - Usually cheaper than locking, but
 - Mistakes are expensive

Third: Lazy Synchronization

- Postpone hard work
- Removing components is tricky
 - Logical removal
 - Mark component to be deleted
 - Physical removal
 - Do what needs to be done

Fourth: Lock-Free Synchronization

- Don't use locks at all
 - Use `compareAndSet()` & relatives ...

Fourth: Lock-Free Synchronization

- Don't use locks at all
 - Use `compareAndSet()` & relatives ...
- Advantages
 - No Scheduler Assumptions/Support
- Disadvantages
 - Complex
 - Sometimes high overhead

Concurrent List

Sequential List Based Set

Add()




Remove()




Sequential List Based Set


add()




remove()



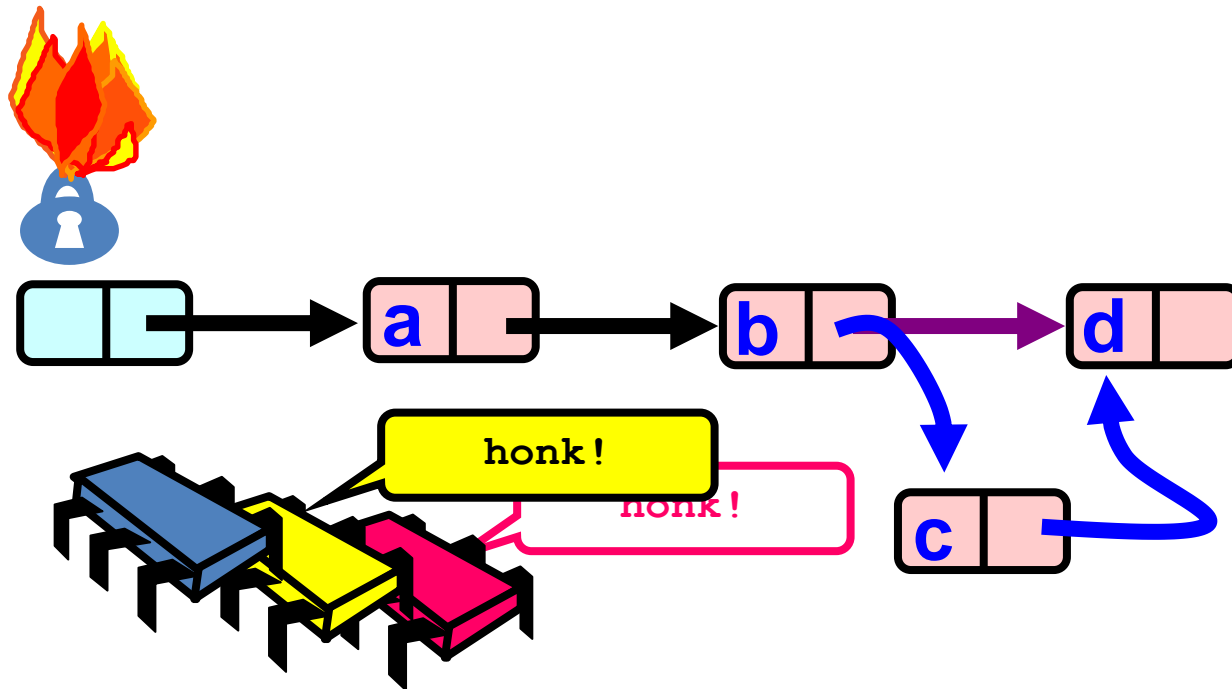
Coarse-Grained Locking



Coarse-Grained Locking



Coarse-Grained Locking



Simple but hotspot + bottleneck

List

- Suppose a list have Nodes, each node have
 - Key value : Customer Name
 - Item Value : Amount
- Node : key, value, next
- Operations frequencies
 - Balance check : 50%
 - Update amount : 49.8
 - Remove account : 0.1%
 - Add account : 0.1%

Fine Gran locking


- Suppose a list have Nodes, each node have
 - Key value : Customer Name
 - Item Value : Amount
- Node : key, value, next
- Operations frequencies
 - Balance check : 50% : Lock not required
 - Update amount : 49.8 : lock of the required node
 - Remove account : 0.1% : lock entire list
 - Add account : 0.1% : lock entire list

Fine Gran locking

- Suppose a list have Nodes, each node have
 - Key value : Customer Name
 - Item Value : Amount
- Node : key, value, next
- Operations frequencies
 - Balance check : 50% : Lock not required
 - Update amount : 49.8 : lock of the required node
 - Remove account : 0.1% : **lock entire list**
 - Add account : 0.1% : **lock entire list**
- If the Add and remove account contribute: 90% of operation, we need to fine grain add and remove : **Refer Book Art of Multiprocessor programming**

Concurrent Hash


Sequential Closed Hash Map



2 Items

$$h(k) = k \bmod 4$$


Add an Item



3 Items

$$h(k) = k \bmod 4$$


Add Another: Collision



4 Items

$$h(k) = k \bmod 4$$


More Collisions



5 Items

$$h(k) = k \bmod 4$$

More Collisions

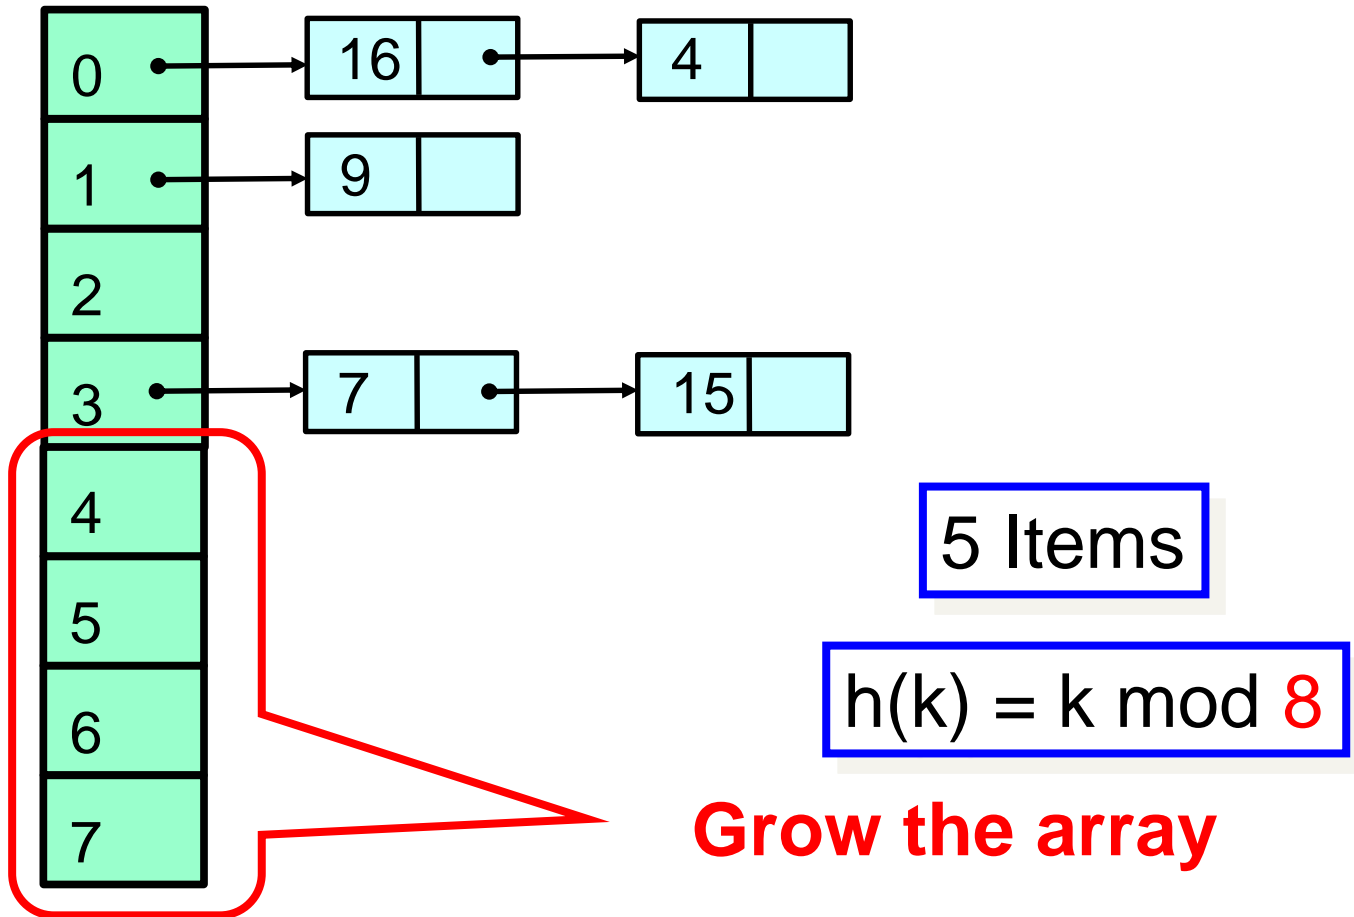


Problem:
buckets getting too long

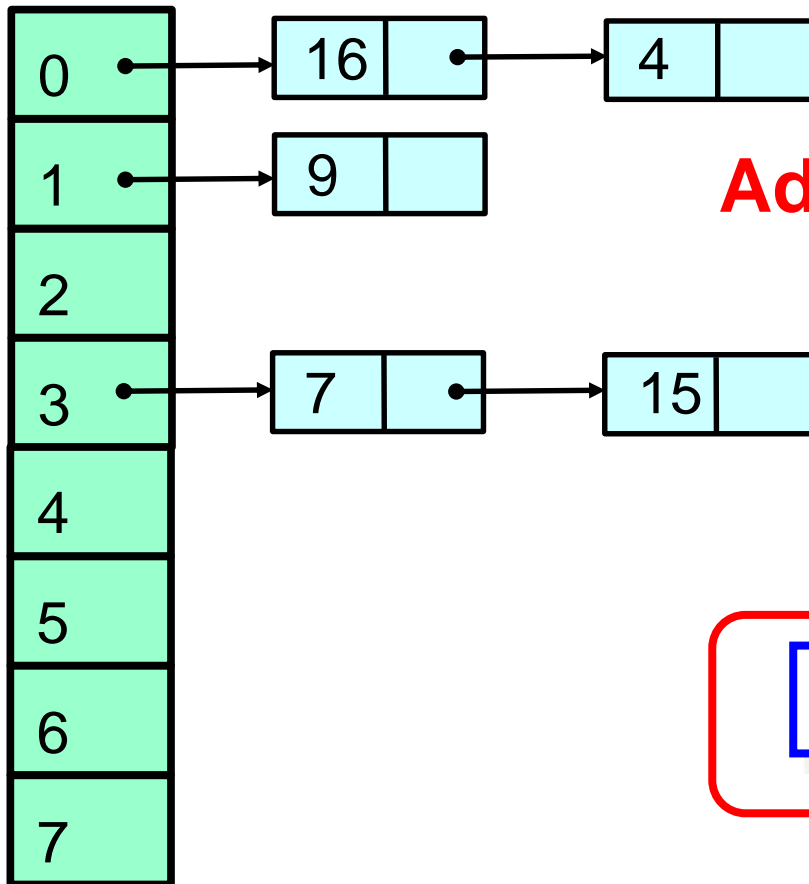
5 Items

$$h(k) = k \bmod 4$$

Resizing



Resizing

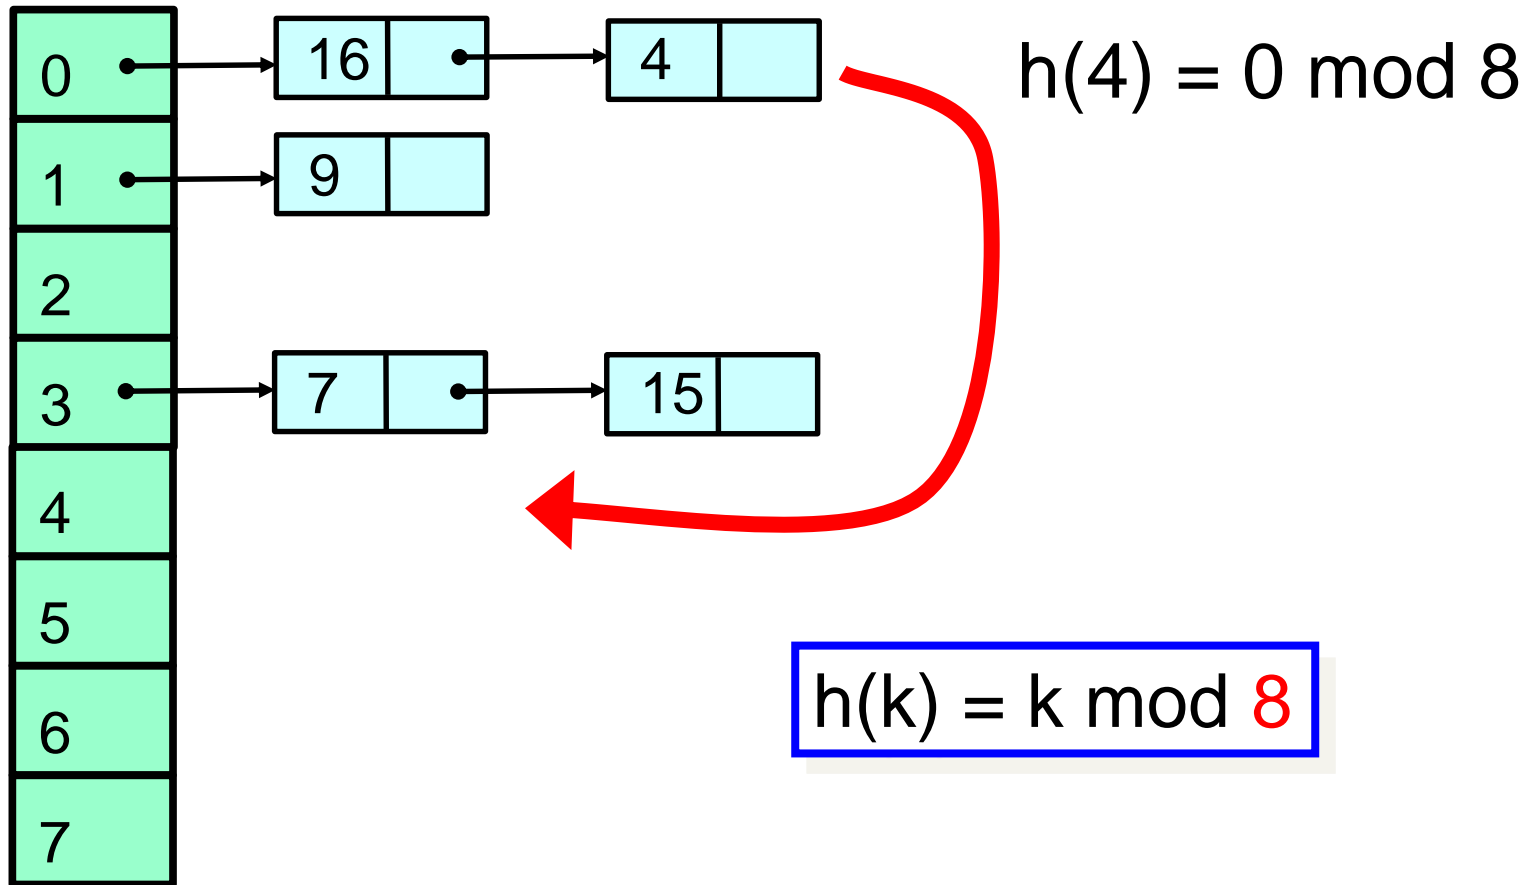


Adjust hash function

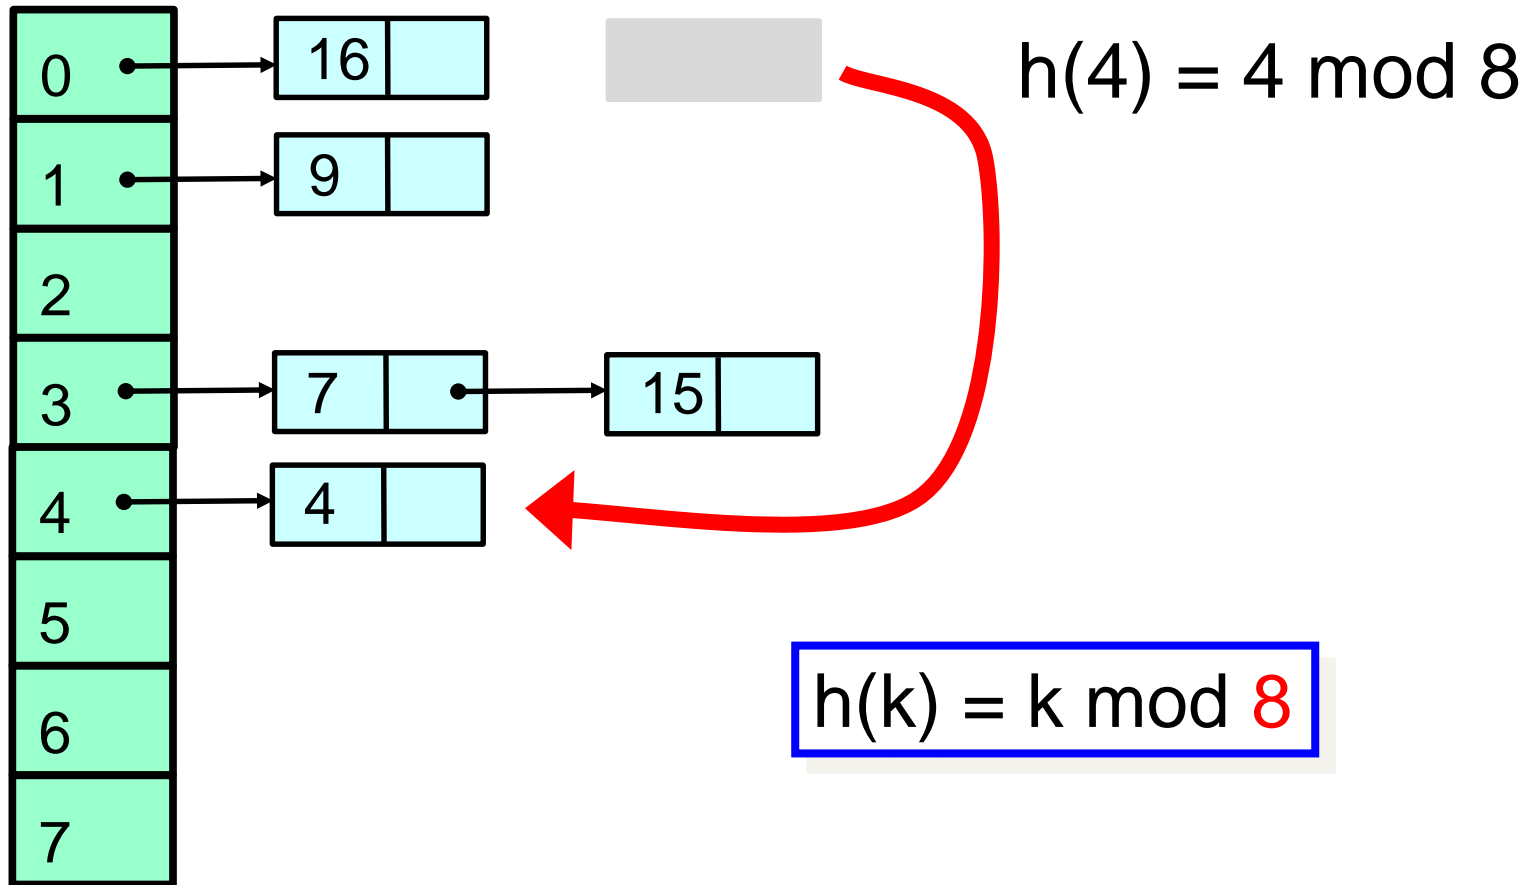
5 Items

$$h(k) = k \bmod 8$$


Resizing



Resizing



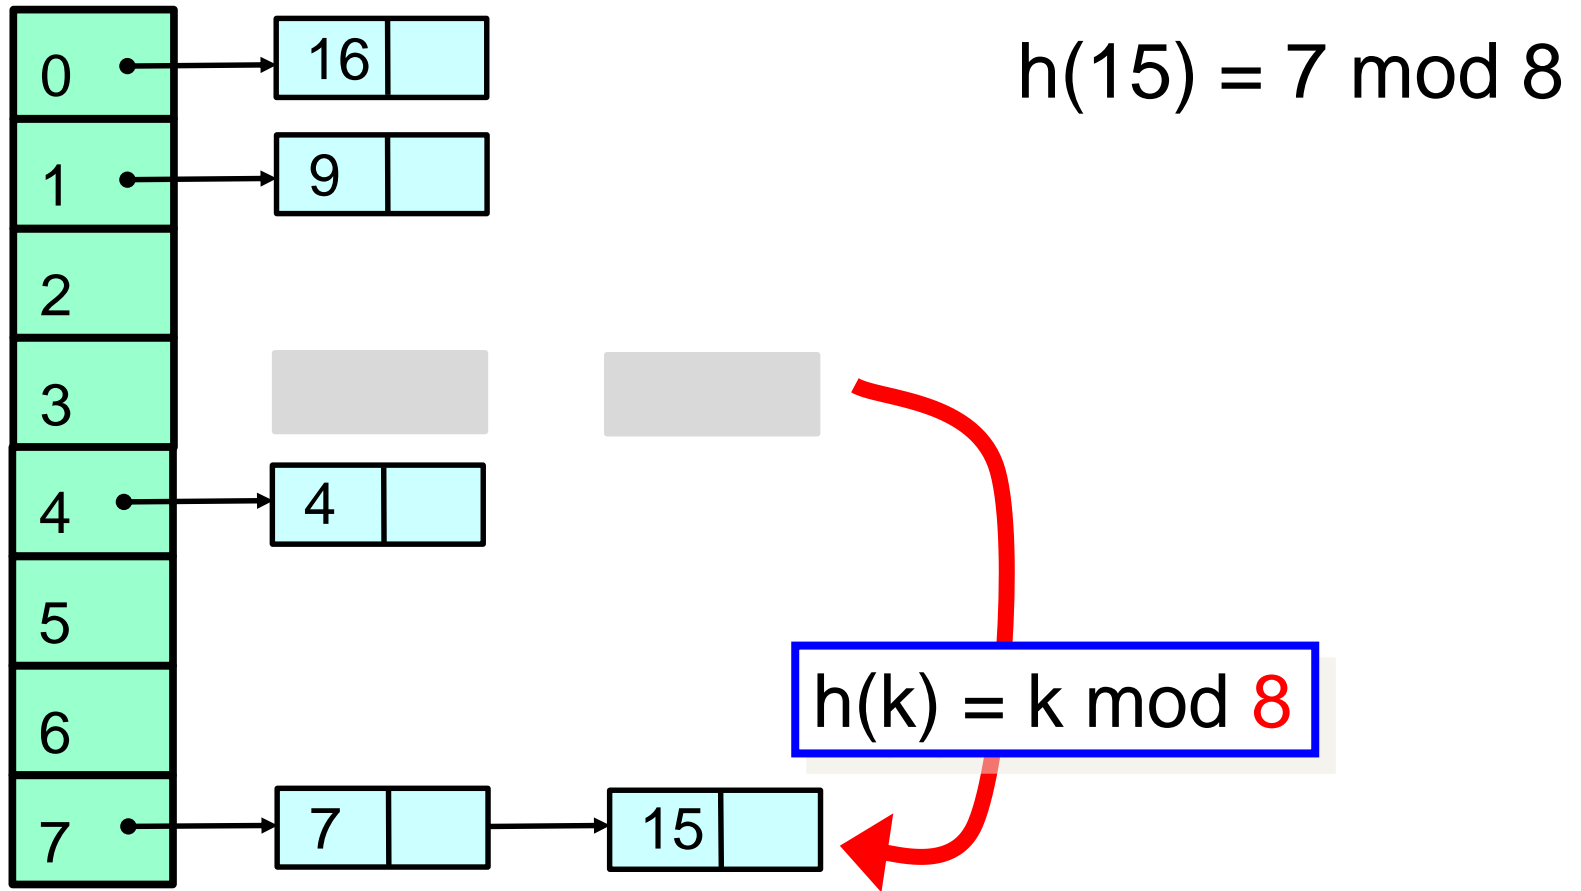
Resizing



$$h(15) = 7 \bmod 8$$

$$h(k) = k \bmod 8$$

Resizing



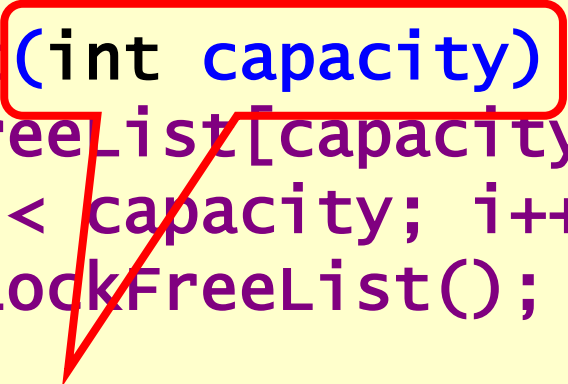
Fields

```
public class SimpleHashSet {  
    protected LockFreeList[] table;  
  
    public SimpleHashSet(int capacity) {  
        table = new LockFreeList[capacity];  
        for (int i = 0; i < capacity; i++)  
            table[i] = new LockFreeList();  
    }  
    ...  
}
```

Array of lock-free lists

Constructor

```
public class SimpleHashSet {  
    protected LockFreeList[] table;  
  
    public SimpleHashSet(int capacity) {  
        table = new LockFreeList[capacity];  
        for (int i = 0; i < capacity; i++)  
            table[i] = new LockFreeList();  
    }  
    ...  
}
```



Initial size

Constructor

```
public class SimpleHashSet {  
    protected LockFreeList[] table;  
  
    public SimpleHashSet(int capacity) {  
        table = new LockFreeList[capacity];  
        for (int i = 0; i < capacity; i++)  
            table[i] = new LockFreeList();  
    }  
    ...
```

Allocate memory

Constructor

```
public class SimpleHashSet {  
    protected LockFreeList[] table;  
  
    public SimpleHashSet(int capacity) {  
        table = new LockFreeList[capacity];  
        for (int i = 0; i < capacity; i++)  
            table[i] = new LockFreeList();  
    }  
    ...  
}
```

Initialization

Add Method

```
public boolean add(Object key) {  
    int hash =  
        key.hashCode() % table.length;  
    return table[hash].add(key);  
}
```

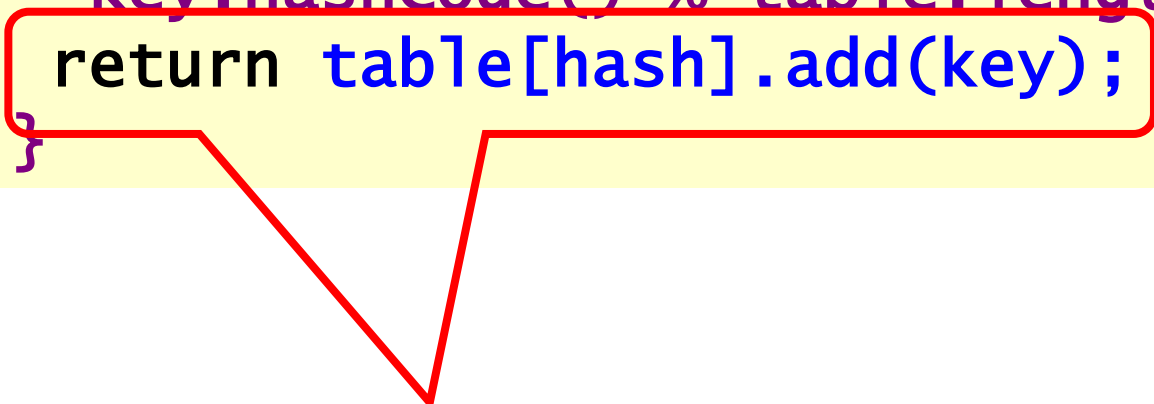
Add Method

```
public boolean add(Object key) {  
    int hash =  
        key.hashCode() % table.length;  
    return table[hash].add(key);  
}
```

**Use object hash code to
pick a bucket**

Add Method

```
public boolean add(Object key) {  
    int hash =  
        key.hashCode() % table.length;  
    return table[hash].add(key);  
}
```



Call bucket's add() method

Is Resizing Necessary?

- Constant-time method calls require
 - Constant-length buckets
 - Table size proportional to set size
 - As set grows, must be able to resize

Set Method Mix

- Typical load
 - **90%** contains() //Read operation
 - **9%** add ()
 - **1%** remove()
- Growing is important
- Shrinking not so much


When to Resize?

- Many reasonable policies. Here's one.
- Pick a threshold on num of items in a bucket
- Global threshold
 - When $\geq \frac{1}{4}$ buckets exceed this value
- Bucket threshold
 - When any bucket exceeds this value

Coarse-Grained Locking

- Good parts
 - Simple
 - Hard to mess up
- Bad parts
 - Sequential bottleneck


Fine-grained Locking




Each lock associated with one bucket

Resize This

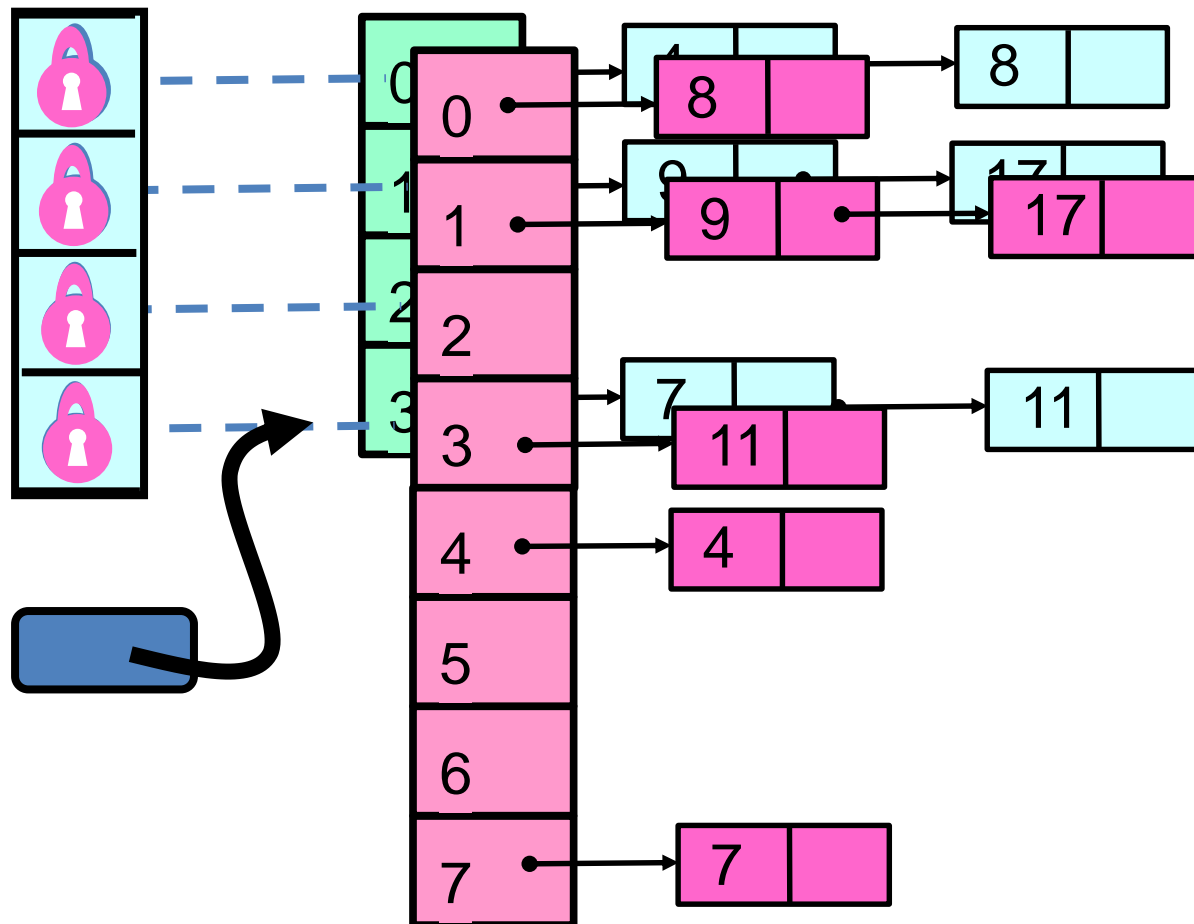
**Make sure table reference didn't change
between resize decision and lock acquisition**



Resize This




Resize This



Striped Locks: each lock now associated with two buckets

Resize This



Observations

- We grow the table, but not locks
 - Resizing lock array is tricky ...
- We use sequential lists
 - If we're locking anyway, why pay?

Assignment II

- Design and implement a Java multithreaded program
- To simulate a bank money transaction system for up to 10^6 users.
- Name of the bank is Guwahati National Bank (GNB)
 - the bank have 10 branches at different location of India and each branch have 10 updaters.
- Initially, 10^4 users will be there for each branch with a random amount of money in their accounts.
- **Every updater of the GNB can be modeled as a separate thread.**

Assignment II

- The updater gets a request
 - To Cash Deposit, Cash withdrawal (withdrawal amount should less than amount money in the account).
 - To Transfer Money from one customer account to another customer account. The source account and destination account may be in different branches of GNB.
 - to add a customer with some initial money in that account, added customer will be in the **updater branch** of GNB.
 - To delete a customer from the system (or close the account of a user of the GNBs from **any branch**).
 - To transfer customer account from **one branch to another branch** of GNB.

Assignment II

- Assume the probability of getting cash deposit, cash withdrawal, money transfer, add a customer, delete a customer, transfer a customer to updater are 0.33, 0.33, 0.33, 0.003, 0.003, and 0.004 respectively.
- Simulate up to 10^6 transactions per updater to test
 - the correctness of your implementation.
 - Report the execution time of your simulation program.

Assignment II

- Suppose information about all the customer accounts of the GNB
 - is maintained by **an array of linked lists**.
 - Each linked list represent (or hold data of a) branch of GNB, so there are 10 linked list and the array of linked list is maintained by a hash data structure.
 - Every customer account number is represented by 10 digits and the first digit of the customer account number identify the branches of GNB.
- You to be used linked list and hash should be
 - thread-safe, and throughput should be high.
 - You are allowed to use any inbuilt data structure, locking protocol, synchronized functions for the same.

Thanks