

CS331 Tutorial 2

Java Thread Programming

A. Sahu

Dept of Comp. Sc. & Engg.


Indian Institute of Technology Guwahati

Outline




- Motivation for Concurrent Programming
 - Earlier Goal : Many thing want to run simultaneously
 - Multi-Tasking using Single Cores: Surfing email at the time of listening songs
 - High Performance as Multi-Cores
 - Multiprocessor programming == > Threading
 - Multi-computer programming == > Multi-processing/process level parallelism
- Threading Methodology
- 8 Rules Designing Mutli-Threaded APPs
- Java Threading Examples

Multi Tasking on a Single Core PC


- Emulating concurrency on single core
- Giving user with experience of multiple task running on single core : **Surfing email at the time of listening songs**
- **Interleaving** is the best way to achieve



Example of Multicore



Tilera64



Why Multicore ?

- Many applications are highly parallel
 - Take benefit of all parallelism (instruction, data and thread)
- Multiprocessors
 - Flexible, programmable, high performance
 - Take benefit of all parallelism (instruction, data and thread)
 - Likely to be cost/power effective solutions



Why Multicore ?

- Multiprocessors are likely to be cost/power effective solutions
 - Share lots of resources
 - *Personal room is costlier than dormitory*
 - *You cannot allocate a Bungalow to each student: it will too costly*
 - *Hostel room with shared facility is sufficient*
 - Need not require very high frequency to run
 - Lots of replication makes easy to manage and cost effective in design

Multicore Difficulties I

- Multiprocessors are likely to be cost/power effective solutions
 - Because it share lots of resources
 - ***Personal room is costlier than dormitory***
 - Sharing resource arise many other problems
 - Critical Sections
 - Lock and Barrier Design
 - Coherence
 - Shared data at all placed should be same
 - Consistency
 - Order should be similar to serial (ROB)
 - One processor Interference others
 - Share efficiently using some policy

Multicore Difficulties II

- Task scheduling in multiprocessors
 - Deterministic task scheduling on multiprocessor with more than 2 processor is NP-Complete problem
 - 4 Tasks (A,B, C and D), 3 Processor
 - {A,B,C,D}{}{} , {A,B,C}{D}{} ,Exponential Number of Solutions

Multicore Difficulties III

- Many applications are highly parallel
 - Take benefit of all parallelism (instruction, data and thread)
 - Most of the coder write sequential code
 - Who will extract parallelism from applications ?
 - There is no successful auto-parallelisation tool till date
 - » Attempts: Cetus, SUIF, SolarisCC

Threading

Traditional View of a Process

- Process = process context + code, data, and stack

Process context

Program context:

Data registers

Condition codes

Stack pointer (SP)

Program counter (PC)


Kernel context:

VM structures (VMem)

Descriptor table

brk pointer

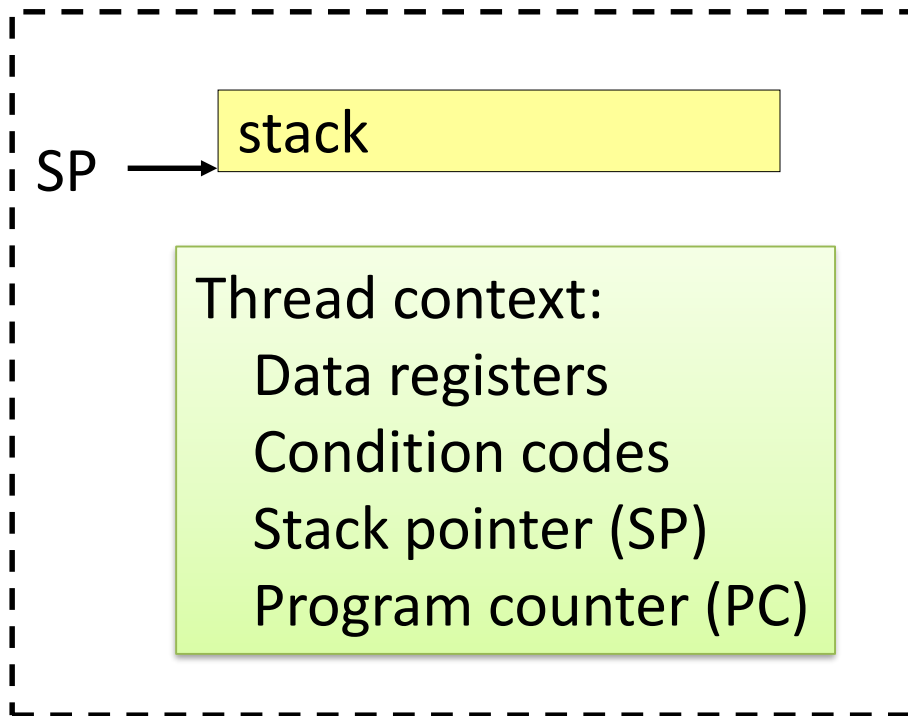
Code, data, and stack



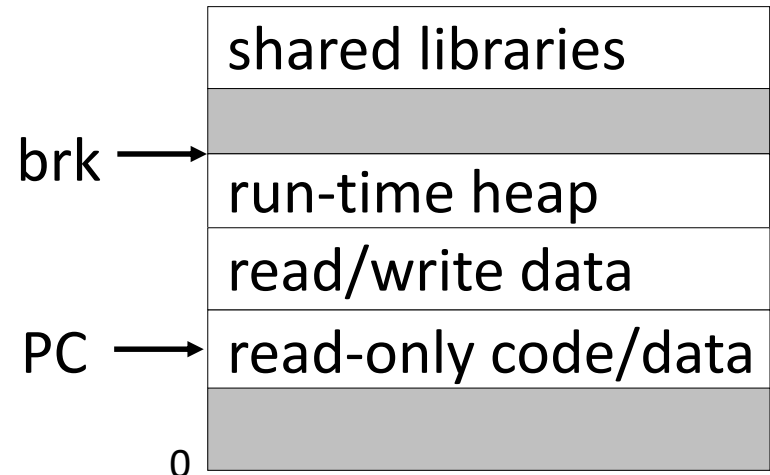
Alternate View of a Process

- Process = thread+ code, data & kernel context

Thread (main thread)



Code and Data



Kernel context:
VM structures
Descriptor table
brk pointer

A Process With Multiple Threads

- Multiple threads can be associated with a process
 - Each thread has its *own* logical control flow (sequence of PC values)
 - Each thread *shares* the same code, data, and kernel context
 - Each thread has its own thread id (TID)

A Process With Multiple Threads

Thread 1 (main thread)

stack 1

Thread 1 context:

Data registers

Condition codes

SP1

PC1

Shared code
and data

shared libraries

run-time heap

read/write data

read-only code/data

0

Kernel context:

VM structures

Descriptor table

brk pointer

Thread 2 (peer
thread)

stack 2

Thread 2 context:

Data registers

Condition codes


SP2

PC2


Logical View of Threads

- Threads associated with a process form a pool of peers
 - Unlike processes, which form a tree hierarchy

Threads associated with
process foo



Process hierarchy



Threads vs. Processes

- How threads and processes are similar
 - Each has its own logical control flow
 - Each can run concurrently
 - Each is context switched

Threads vs. Processes

- How threads and processes are different
 - Threads share code and data, processes (typically) do not
 - Threads are somewhat less expensive than processes
 - Process control (creating and reaping) is twice as expensive as thread control
 - Linux/Pentium III numbers:
 - ~20K cycles to create and reap a process
 - ~10K cycles to create and reap a thread

Threading Language and Support

- Pthread: POSIX thread
 - Popular, Initial and Basic one
- Improved Constructs for threading
 - c++ thread : available in c++11, c++14
 - **Java thread : very good memory model**
 - **Atomic function, Mutex**
- Thread Pooling and higher level management
 - OpenMP (loop based)
 - Cilk (dynamic DAG based)

Threading Methodology

Threading Methodology

- **Does not recommend going straight to concurrency!**
- First produce a tested single-threaded program
 - Use reqs./ design/ implement /test/ tune/ maintenance steps
- Then to create a concurrent system from the former, do
 1. Analysis: Find computations that are independent of each other
 1. AND take up a large amount of serial execution time (80/20 rule)
 2. Design and Implement: straightforward Test for Correctness: Verify that concurrent code produces correct output
 3. Tune for performance: once correct, find ways to speed up

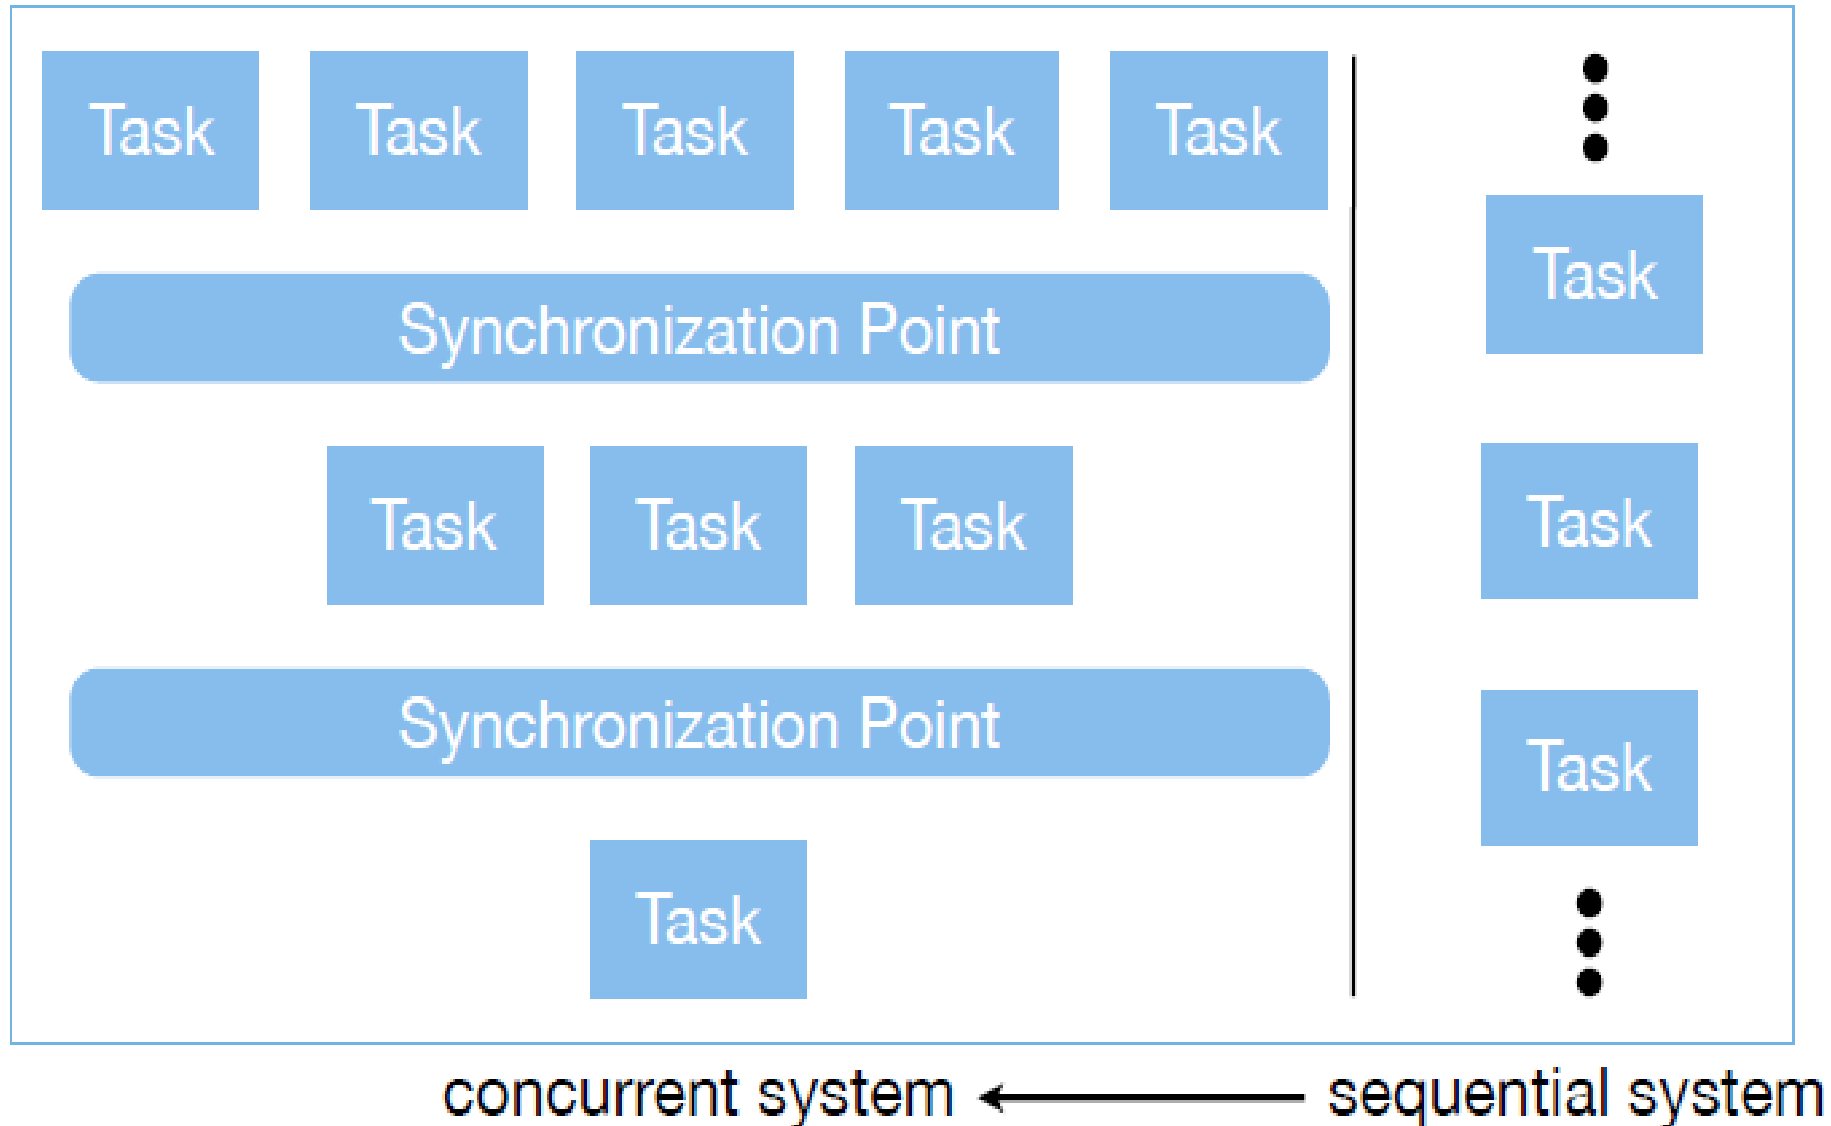
Performance Tuning

- Tuning threaded code typically involves
 - identifying sources of contention on locks (synchronization)
 - identifying work imbalances across threads
 - reducing overhead
- Testing and Tuning
 - Whenever you tune a threaded program, you must test it again for correctness
- Going back further
 - if you are unable to tune system performance,
 - you may have to re-design and re-implement

Design Models

- Two primary design models for concurrent algorithms
- Task Decomposition
 - identify tasks (computations) that can occur in any order
 - assign such tasks to threads and run concurrently
- Data Decomposition
 - program has large data structures where individual data elements can largely be calculated independently
 - data decomposition implies task decomposition in these cases

Task Decomposition



Eight rules of Designing multithreaded APPS

1/8 Rules: Designing multithreaded APPS

- **Identify Truly Independent Computations**
- If you can't identify (in a single threaded application) computations that can be done in parallel, you're out of luck
- Some situations that indeed can't be made parallel

2/8 Rules: Designing multithreaded APPS

- **Implement Concurrency at the Highest Level Possible**
- When discussing “What’s Not Parallel” a common refrain was “you can’t make this parallel,
 - So see if its part of a larger computation that CAN be made parallel”
- This is such good advice, it was promoted to being a guideline!
 - Two approaches: bottom up, top down

2/8 Rules: Bottom UP

- Our methodology says to create a concurrent program
 - Start with a tuned, single-threaded program and
 - Use a profiler to find out where it spends most of its time
- In the bottom-up approach, you start at those “hot spots” and work up; typically, a hotspot will be a loop of some sort
 - See if you can thread the loop
 - If not, move up the call chain, looking for the next loop and see if it can be made parallel...
 - If so, still look up the call chain for other opportunities, first.
 - Why? Granularity! You want coarse-grained tasks for your thread

2/8 Rules: Top Down

- With knowledge of the location of the hot spot
- Start by looking at the whole application and see if there are parallelization opportunities on the large-scale structure that contains the hot spot
 - if so, you've probably found a nice coarse-grained task to assign to your threads
 - If not, move lower in the code towards the hot spot, looking for the first opportunity to make the code concurrent

3/8 Rules: Designing multithreaded APPS

- **Plan Early for Scalability**
- The number of cores will keep increasing
- You should design your system to take advantage of more cores as they become available
 - Make the number of cores an input variable and design from there
- In particular, designing systems via data decomposition techniques will provide more scalable systems
 - humans are always finding more data to process!
- More data, more tasks; if more cores arrive, you're ready

4/8 Rules: Designing multithreaded APPS

- **Make use of Thread-Safe Libraries Wherever Possible**
- First, software reuse!
 - Don't fall prey to Not Invented Here Syndrome
 - if code already exists to do what you need, use it!
- Second, more libraries are becoming multithread aware
 - That is, they are being built to perform operations concurrently
- Third, if you make use of libraries, ensure they are thread-safe; if not, you'll need to synchronize calls to the library
 - Global variables hiding in the library may prevent even this, if the code is not reentrant ; if so, you may need to abandon it

5/8 Rules: Designing multithreaded APPS

- **Use the Right Threading Model**
- Avoid the use of explicit threads if you can get away with it
- They are hard to get right
- Look at libraries that abstract away the need for explicit threads
 - OpenMP, Cilk and Intel Threading Building Blocks
 - Scala's agent model, Go's goroutines and Clojure's concurrency primitives
- All of these models hide explicit threads from the programmer Right Threading Model

6/8 Rules: Designing multithreaded APPS

- **Never Assume a Particular Order of Execution**
- With multiple threads, as we've seen, the scheduling of atomic statements is nondeterministic
- If you care about the ordering of one thread's execution with respect to another, you have to impose synchronization
- But, to **get the best performance**, you want to **avoid synchronization** as much as possible
- In particular, you want high granularity tasks that don't require synchronization
 - This allows your cores to run as fast as possible on each task they're given

7/8 Rules: Designing multithreaded APPS

- **Use Thread-Local Storage Whenever Possible or Associate Locks with specific data**
- Related to Rule 6; the more your threads can use thread-local storage, the less you will need synchronization
- Otherwise, associate a single lock with a single data item
 - in which a data item might be a huge data structure
- This makes it easier for the developer to understand the system;
 - “if I need to update data item A, then I need to acquire lock A first”

8/8 Rules: Designing multithreaded APPS

- **Dare to Change the Algorithm for a Better Chance of Concurrency**
- Sometimes a tuned, single-threaded program makes use of an algorithm which is not amenable to parallelization
- They might have picked that algorithm for performance reasons
 - Strassen's Algorithm $O(n^{2.81})$ vs. the triple-nested loop algorithm to perform matrix multiplication $O(n^3)$
- Change the algorithm used by the single-threaded program to see if you can then make that new algorithm concurrent
 - BUT: when measuring speedup, compare to the original!!

Thread Programing In Java

What Is a Java Thread?

- A thread is actually a lightweight process.
- Java provides built-in support for **multithreaded programming**.
- A multithreaded program contains two or more parts that can run **concurrently**.
 - Each part of such a program is called a thread and
 - each thread defines a separate path of the execution.
- Thus, multithreading is a specialized form of multitasking.

Java Thread Model

- The Java run-time system depends on threads for many things.
- Threads reduce inefficiency by preventing the waste of CPU cycles.
- Threads exist in several states:
 - **New** - When we create an instance of Thread class, a thread is in a new state.
 - **Running** - The Java thread is in running state.
 - **Suspended** - A running thread can be **suspended**, which temporarily suspends its activity. A suspended thread can then be resumed, allowing it to pick up where it left off.
 - **Blocked** - A Java thread can be blocked when waiting for a resource.
 - **Terminated** - A thread can be terminated, which halts its execution immediately at any given time.

Multithreading in Java: Thread Class and Runnable Interface

- Java's multithreading system is built upon the Thread class
- Its methods, and its companion interface, **Runnable**.
- To create a new thread, your program will
 - either extend **Thread**
 - or **implement** the **Runnable** interface.

Thread class methods to manage threads

<code>getName()</code>	Obtain thread's name
<code>getPriority()</code>	Obtain thread's priority
<code>isAlive()</code>	Determine if a thread is still running
<code>join()</code>	Wait for a thread to terminate
<code>run()</code>	Entry point for the thread
<code>sleep()</code>	Suspend a thread for a period of time
<code>start()</code>	Start a thread by calling its run method

Main Java Thread

- Thread and Runnable interface to create and manage threads, beginning with the **main java thread**.
- **Why Is Main Thread So Important?**
 - Because it affects the other 'child' threads.
 - Because it performs various shutdown actions.
 - Because it's created automatically when your program is started.

How to Create a Java Thread

- Java lets you create a thread one of two ways:
 - By **implementing** the **Runnable** interface.
 - By **extending** the **Thread**.
- **Runnable Interface**
 - The easiest way to create a thread is to create a class that implements the **Runnable** interface.
 - To implement Runnable interface, a class need only implement a single method called run()
- **Extending Java Thread**
 - Override the run() method and then to create an instance of that class.
 - The run() method is what is executed by the thread after you call start()

Create Java Thread: Runnable Interface

```
public class MyThIR implements Runnable{  
    public void run() {  
        System.out.println("Child  
        thread is running...");  
    }  
}  
  
public static void main(String args[]) {  
    MyThIR m1=new MyThIR();  
    Thread t1 =new Thread(m1);  
    t1.start();  
}
```

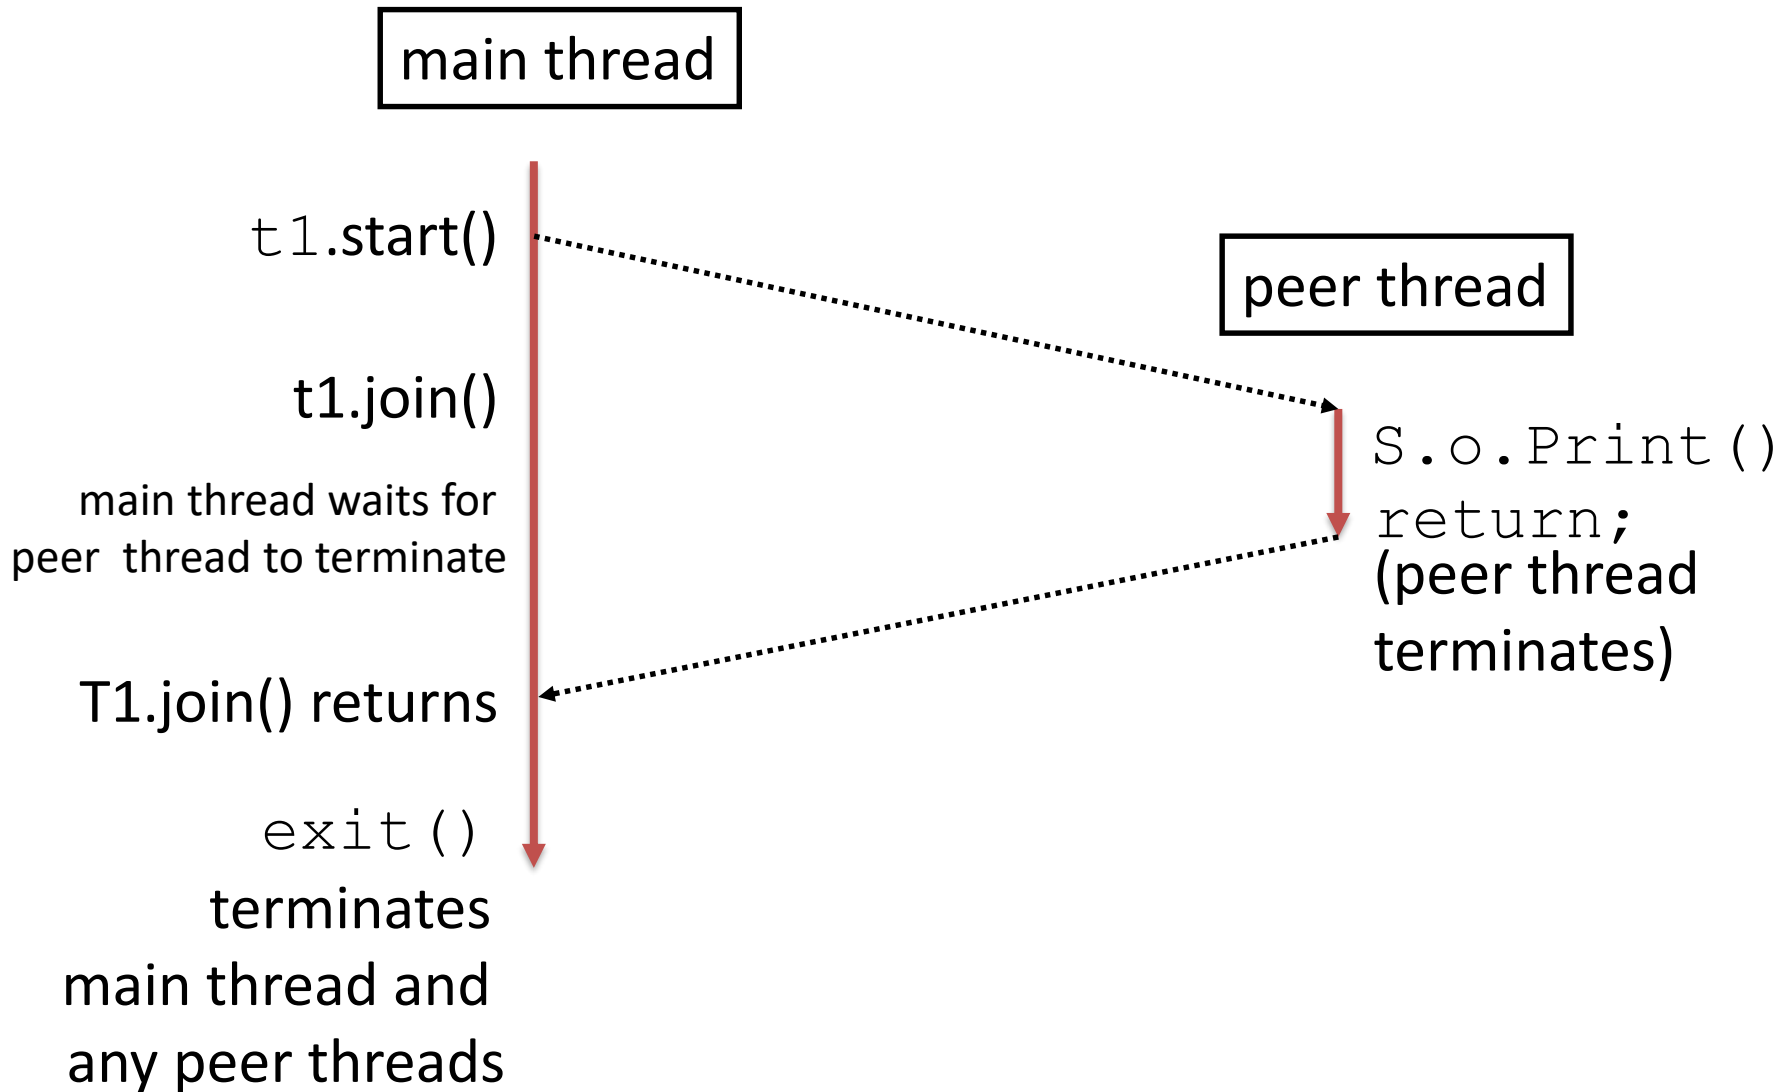
Create Java Thread: Extending Java Thread

```
public class MyTh_EJT extends Thread{  
    public void run() {  
        System.out.println("Child  
        thread is running...");  
    }  
}  
  
public static void main(String args[]) {  
    MyTh_EJT t1 = new MyTh_EJT();  
    t1.start();  
}
```

Create Java Thread: Extending Java Thread

```
public class MyTh_EJT extends Thread{  
    public void run() {  
        System.out.println("Child  
        thread is running...");  
        return;  
    }  
}  
  
public static void main(String args[]) {  
    MyTh_EJT t1 = new MyTh_EJT();  
    t1.start();  
    t1.join();  
}
```

Execution of Threaded “child thread”



VectorSum Serial

```
int A[VSize], B[VSize], C[VSize];  
void VectorSumSerial() {  
    for( int j=0; j<SIZE; j++)  
        A[j]=B[j]+C[j];  
}
```

Suppose Size=1000

0-249	250-499	500-749	750-999
-------	---------	---------	---------

T1

T2

T3

T4

VectorSum Serial

```
int A[VSize], B[VSize], C[VSize];

void VectorSumSerial() {
    for( int j=0; j<SIZE; j++)
        A[j]=B[j]+C[j];
}
```

- Independent
- Divide work into equal for each thread
- Work per thread: $\text{Size}/\text{numThread}$

VectorSum Parallel

```
void DoVectorSum(int TID) {  
    int j, SzPerthrd, LB, UB;  
    SzPerthrd=(VSize/NUM_THREADS);  
    LB= SzPerthrd*TID;UB=LB+SzPerthrd;  
  
    for (j=LB; j<UB; j++)  
        A[j]=B[j]+C[j];  
}
```


VectorSum Parallel : Java Threads

```
public class MyTh_EJT extends Thread{
    public void run(int TID) {
        DoVectorSum(TID)
        return;
    }
}

public static void main(String args[]) {
    int NumOfTasks = Integer.parseInt(args[0]);
    for ( int i=0; i < NumOfTasks; i++) {
        MyTh_EJT t1=new MyTh_EJT();
        t1.start(i);
    }
}
```



VectorSum Parallel : Java Threads

Complete java program will be
uploaded to MS Team

Matrix multiply and threaded matrix multiply

- Matrix multiply: $C = A \times B$

$$C[i, j] = \sum_{k=1}^N A[i, k] \times B[k, j]$$



Matrix multiply and threaded matrix multiply

- Matrix multiply: $C = A \times B$

$$C[i, j] = \sum_{k=1}^N A[i, k] \times B[k, j]$$

- Divide the whole rows to T chunks
 - Each chunk contains : N/T rows, Assume $N\%T=0$

Four ways to implement a synchronized counter in Java

- **Suppose there is a Shared counter and every threads are attempting to manipulate**
 - 1. Synchronized Block**
 - 2. Atomic Variable**
 - 3. Concurrent Lock**
 - 4. Semaphore**

Simple Java Ctr: without Lock

```
import java.util.concurrent.ExecutorService;
import java.util.concurrent.Executors;

class Counter implements Runnable {
    private static int counter = 0;
    private static final int limit = 1000;
    private static final int threadPoolSize = 5;

    public static void main(String[] args) {
        ExecutorService ES = Executors.newFixedThreadPool(threadPoolSize);
        for (int i = 0; i < threadPoolSize; i++) { ES.submit(new Counter()); }
        ES.shutdown();
    }
}
```

Simple Java Ctr: without Lock

```
@Override
public void run() {
    while (counter < limit) {
        increaseCounter();
    }
}

private void increaseCounter() {
    System.out.println(Thread.currentThread().getName() + " : " + counter);
    counter++;
}
}
```

Simple Java Ctr: with Sync Lock

```
import java.util.concurrent.ExecutorService;
import java.util.concurrent.Executors;

class Counter implements Runnable {
    private static int counter = 0;
    private static final int limit = 1000;
    private static final int threadPoolSize = 5;
    private static final Object lock = new Object();
    public static void main(String[] args) {
        ExecutorService ES = Executors.newFixedThreadPool(threadPoolSize);
        for (int i = 0; i < threadPoolSize; i++) { ES.submit(new Counter()); }
        ES.shutdown();
    }
}
```


Simple Java Ctr: without Lock

```
@Override
public void run() {
    while (counter < limit) {
        increaseCounter();
    }
}

private void increaseCounter() {
    synchronized (lock) {
        System.out.println(Thread.currentThread().getName() + " : " + counter);
        counter++;
    }
}
}
```

Simple Java Ctr: with atomic Int

```
import java.util.concurrent.ExecutorService;
import java.util.concurrent.Executors;
import java.util.concurrent.atomic.AtomicInteger;
class Counter implements Runnable {
    private static AtomicInteger counter;
    private static final int limit = 1000;
    private static final int threadPoolSize = 5;
    public static void main(String[] args) {
        ExecutorService ES = Executors.newFixedThreadPool(threadPoolSize);
        for (int i = 0; i < threadPoolSize; i++) { ES.submit(new Counter()); }
        ES.shutdown();
    }
}
```

Simple Java Ctr: atomic int

```
@Override
public void run() {
    while (counter < limit) {
        increaseCounter();
    }
}

private void increaseCounter() {
    System.out.println(Thread.currentThread().getName() + " : " + counter);
    counter.getAndIncrement();
}
}
```

Thanks